

Joining Data Sets using Std. Python Libraries

Problem

As a stakeholder, I want to receive a report about non-blocked transactions in terms of cumulative transactions amount and total number of unique active users broken down by transaction category.

Proposed solution

Unfortunately the query planner of our database can not optimize the query well enough in order to get the results. It is proposed to develop a python application to join two tables, `users` and `transactions`:

```
CREATE TABLE transactions
(
    transaction_id      UUID,
    date                DATE,
    user_id             UUID,
    is_blocked          BOOL,
    transaction_amount  INTEGER,
    transaction_category_id INTEGER
);

CREATE TABLE users
(
    user_id      UUID,
    is_active    BOOLEAN
);
```

The application shall implement the business logic described in the query below.

Query

```
SELECT t.transaction_category_id,
       SUM(t.transaction_amount) AS sum_amount,
       COUNT(DISTINCT t.user_id) AS num_users
FROM transactions t
      JOIN users u USING (user_id)
WHERE t.is_blocked = False
      AND u.is_active = 1
GROUP BY t.transaction_category_id
ORDER BY sum_amount DESC;
```

Acceptance Criteria

- Python application is developed to perform required calculations and print the result to *stdout*.
- The application satisfies the following success criteria:

- [Correctness](#)
- [Code quality](#)
- [Efficiency](#)
- Only the standard python libraries are allowed to be used.
- The solution does not concern SQL parsing.
- The solution is not generic, it does only concern the described problem.

Correctness check

Data quality assurance is performed according to the following test scenario.

GIVEN

- The inputs `users.csv` and `transactions.csv` generated using [the script](#);

WHEN

- The logic application is executed end-to-end;

THEN

- The program output matches the reference generated by postgres.

Code quality

Linting applied:

- black
- flak8
- mypy
- isort

Efficiency

Metrics:

- Execution time
- Memory allocation

References:

- <https://wiki.python.org/moin/TimeComplexity>
- <https://realpython.com/sorting-algorithms-python/>

Solution

The solution is provided in the [solution](#) directory.

How to run

Requirements:

- [docker](#) ~> 20.10

- [gnuMake](#)

Note: the docker compose v2 used, i.e. `docker compose` instead of `docker-compose` command.

Commands

Run to see available commands:

```
make help
```

First, set the environment:

```
make setup
```

Second, run the unit tests and the [data quality assessment](#):

```
make tests
```

To run the application, execute the following command:

```
make run
```

Note: the application uses the csv files generated at the setup step.

Consider the following options to use custom `users.csv` and `transactions.csv` files as the input data:

- move the files to `./fixture/tests` directory and run the application;
- place the csv files in a single directory and execute the command:

```
make run BASE_DIR=##path/to/users/and/transactions/csv##
```

Run to perform application profiling on pre-generated data:

```
make profiling
```

Product Questions

- Does input's validation required?
 - The solution is being delivered under the assumption of positive answer.
The reason: data quality for analytics is more critical than performance.

Tech Questions and Decisions

- How to store `user_id` and `transaction_id` in memory?
 - UUID vs. str: what leads to higher memory allocation?
The answer: str required more memory to store an object:

```
In [1]: uid = "9f709688-326d-4834-8075-1a477d590af7"

In [2]: uid.__sizeof__()
Out[2]: 85

In [3]: from uuid import UUID

In [4]: uid_uuid = UUID(uid)

In [5]: uid_uuid.__sizeof__()
Out[5]: 40
```

- Would use of `dataclass` lead to higher memory allocation?

The answer: no, en contraire; on top, it improves code quality:

```
In [1]: from uuid import UUID
In [2]: Transaction = tuple[UUID, UUID, int, int]

In [3]: from dataclasses import dataclass

In [4]: @dataclass
...: class Transaction1:
...:     transaction_id: UUID
...:     user_id: UUID
...:     transaction_amount: int
...:     transaction_category_id: int
...:
In [5]: t0 = Transaction((UUID("9f709688-326d-4834-8075-1a477d590af7"),
UUID("999eb541-c1a0-4888-aeb6-92773fc60e69"), 1, 1))

In [6]: t1 = Transaction1(UUID("9f709688-326d-4834-8075-1a477d590af7"),
UUID("999eb541-c1a0-4888-aeb6-92773fc60e69"), 1, 1)

In [7]: t0.__sizeof__()
Out[7]: 56

In [8]: t1.__sizeof__()
Out[8]: 32
```

- Which sorting algorithm would suffice the cardinality of the problem?

- Costs-benefit tradeoff: delivery effectiveness vs. technical efficiency, i.e. development complexity vs. time-complexity gain.

Hypothesis: the standard `sort` would be as efficient as it is feasibly possible because the data cardinality, or the number of sorted array's elements does not exceed a dozen. The hypothesis is based on personal user experience, provided fixtures generation [script](#) and [the article](#) illustrating different sorting algorithms implementation in python.

- When reading file, is `csv library` more efficient than line-by-line reading with `open`?

Execution path

1. Store the ID of *active* users in memory as a `Set`:

- Benefits:
 - Minimisation of memory allocation:
 - Only the smaller dataset is stored in memory in full;
 - Only unique users are stored in memory.
- Requirements:
 - `WHERE` clause condition to be applied in-flight: filtering for `is_active` when reading the file line-by-line.
 - The array of `user_id` is fully stored in memory to realise the join condition.

2. Extract *relevant* attributes of non-blocked transactions done by *active* users:

- Benefits:
 - Minimisation of memory allocation
- Requirements:
 - `WHERE` clause condition to be applied in-flight: filtering for `is_blocked` when reading the file line-by-line.
 - `JOIN` clause condition to be applied in-flight: filtering for `user_id` to be in the set from the step 1.
 - Select `transaction_category_id`, and `transaction_amount` only to deliver required data.

3. Store the map of `transaction_category_id` to cumulative `transaction_amount` and array of *unique* `user_id` in memory as a `Dict`.

- Benefits:
 - Minimisation of memory allocation for `transaction_amount`.
 - Minimisation of memory allocation by preserving unique `user_id` only.
- Requirements:
 - The equivalent of the `query` operation `SUM(t.transaction_amount)` to be applied on the fly.
 - Uniqueness of `user_id` associated with a given category is guaranteed by design of the `Set` data type.
- Limitations:
 - Array of `user_id` would have to be preserved to keep state of users mapped to a given transaction category to execute the equivalent of the `query` operation `COUNT(DISTINCT user_id)`.

4. Calculate the number of unique active users associated with the transaction category.

5. Sort by the total transaction amount.

6. Output.

Solution Limitations

The whole set of active users identifiers has to be stored in memory. What if it does not fit?

It is the problem for large datasets which could be typically resolve either by vertical, or horizontal scaling of the computation unit.

Vertical scaling is the straightforward approach: increase the amount of memory, so it fits the amount of data.

Horizontal scaling is achieved by parallel execution of computations following the "map-reduce" logic:

- *Map*:
 - The data are distributed across a cluster of computation nodes according to distribution criteria, e.g. using round-robin, or hashing algorithm.
 - The computation for a certain data set is performed on a certain node.
- *Reduce*:
 - The Map's result is collected on a single node and aggregated to get the final result.

Apart from networking and orchestration overhead, the *reduce* operation could hit the resource bottleneck in line with the initial question. In such case, the "map-reduce" process could be repeated, or the resources quota for the "reduce node" could be raised.

Performance Analysis

The section touches upon the logic performance.

The benchmarking was performed on the data generated using the [script](#):

- users.csv with 1 Million rows
- transactions.csv with 100 Million rows

Logic	Elapsed Time [sec.]	RAM uplift [Mb]	CPU max [% of .5 unit]
Reference	70.184	~ 50	< 60
Solution	434.595	~ 250	< 60

Note: the benchmark is based on a single run on a MacBook Pro with Apple M1 Pro and 16Gb of RAM. It shall only be considered as a qualitative illustration rather than quantitative thorough comparison taking statical significance into account.

Reference

Postgres with the buffer setting of 128 kB (minimal possible value) is used as the reference. Its performance is illustrated as following:

```
EXPLAIN ANALYSE INSERT INTO benchmark.result
SELECT t.transaction_category_id,
       SUM(t.transaction_amount) AS sum_amount,
       COUNT(DISTINCT t.user_id) AS num_users
```

```

FROM benchmark.transactions t
      JOIN benchmark.users u USING (user_id)
WHERE NOT t.is_blocked
      AND u.is_active
GROUP BY t.transaction_category_id
ORDER BY sum_amount DESC
;
QUERY
PLAN

```

```

-----
-----
Insert on result (cost=2214662.55..2214662.74 rows=11 width=12) (actual
time=70110.004..70110.004 rows=0 loops=1)
  -> Subquery Scan on "*SELECT*" (cost=2214662.55..2214662.74 rows=11
width=12) (actual time=70109.873..70109.875 rows=11 loops=1)
    -> Sort (cost=2214662.55..2214662.58 rows=11 width=20) (actual
time=70109.861..70109.861 rows=11 loops=1)
      Sort Key: (sum(t.transaction_amount)) DESC
      Sort Method: quicksort Memory: 25kB
    -> GroupAggregate (cost=2209223.12..2214662.36 rows=11
width=20) (actual time=69394.877..70109.832 rows=11 loops=1)
      Group Key: t.transaction_category_id
      -> Sort (cost=2209223.12..2210582.90 rows=543913
width=24) (actual time=69307.383..69539.604 rows=899728 loops=1)
        Sort Key: t.transaction_category_id
        Sort Method: external merge Disk: 29872kB
      -> Merge Join (cost=2131723.26..2146252.70
rows=543913 width=24) (actual time=67670.559..68884.249 rows=899728
loops=1)
        Merge Cond: (t.user_id = u.user_id)
        -> Sort (cost=2048810.31..2051193.65
rows=953335 width=24) (actual time=66815.998..67139.747 rows=999824
loops=1)
          Sort Key: t.user_id
          Sort Method: external merge Disk:
33192kB
        -> Seq Scan on transactions t
(cost=0.00..1934580.64 rows=953335 width=24) (actual time=2.073..65180.612
rows=999824 loops=1)
          Filter: (NOT is_blocked)
          Rows Removed by Filter:
99000176
        -> Materialize (cost=82912.95..85795.37
rows=576485 width=16) (actual time=854.553..1347.073 rows=1231088 loops=1)
          -> Sort (cost=82912.95..84354.16
rows=576485 width=16) (actual time=854.548..1218.244 rows=900081 loops=1)
            Sort Key: u.user_id
            Sort Method: external merge
            Disk: 22824kB
          -> Seq Scan on users u
(cost=0.00..17899.70 rows=576485 width=16) (actual time=0.135..256.031
rows=900081 loops=1)
            Filter: is_active

```

Rows Removed by Filter:

99919
Planning time: 1.170 ms
Execution time: 70183.707 ms
(27 rows)

Results:

```
SELECT * FROM benchmark.result;
transaction_category_id | sum_amount | num_users
-----+-----+-----
5 | 411126340 | 78431
8 | 410552270 | 78442
9 | 410413764 | 78567
0 | 410069189 | 78288
3 | 409259459 | 78225
6 | 408855294 | 78056
10 | 408843738 | 78339
2 | 408564886 | 78055
7 | 408210562 | 77881
4 | 407689371 | 77753
1 | 407210378 | 77939

(11 rows)
```

Database read from cache illustration:


```
SELECT heap_blks_read, heap_blks_hit from pg_statio_user_tables where
relname='transactions' and schemaname='benchmark';
heap_blks_read | heap_blks_hit
-----+-----
1899420 | 1868898
(1 row)

SELECT heap_blks_read, heap_blks_hit from pg_statio_user_tables where
relname='users' and schemaname='benchmark';
heap_blks_read | heap_blks_hit
-----+-----
19113 | 12735
```

The resources consumption assessed using **docker stats**:

- CPU: up to 50% of 0.5 CPU
- RAM: up to 150Mb from 100Mb

Solution




```
make run BASE_DIR=${PWD}/fixtures/benchmark
2022-11-06T22:07:55.006 [INFO] elapsed time: 434.595 sec.
transaction_category_id,sum_amount,num_users
5,411126340,78431
8,410552270,78442
9,410413764,78567
0,410069189,78288
3,409259459,78225
6,408855294,78056
10,408843738,78339
2,408564886,78055
7,408210562,77881
4,407689371,77753
1,407210378,77939
```

The resources consumption assessed using `docker stats`:

- CPU: up to 50% of 0.5 CPU
- RAM: up to 270Mb from <10Mb